

# Machine Learning: Course Project

*Jeremy Peters*

*February 6, 2018*

## Executive Summary

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. In a study, Six young health participants were asked to perform one set of 10 repetitions of the Unilateral Dumbbell Biceps Curl in five different fashions: exactly according to the specification (Class A), throwing the elbows to the front (Class B), lifting the dumbbell only halfway (Class C), lowering the dumbbell only halfway (Class D) and throwing the hips to the front (Class E). Only Class A corresponds to correct performance. The objective of this project is to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants to build a machine learning algorithm to predict the manner/class type in which an exercise was completed. More information about the study and data set can be found in the section on the Weight Lifting Exercise Dataset at the following URL: <http://groupware.les.inf.puc-rio.br/har>. The training data for this project was download from the following URL: <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>. The test data for this project was download from the following URL: <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

## Exploratory Data Analysis

- Read the Training and Testing CSV files in table format, specify types of missing values (NA, empty strings and div0), and create data frames
- Display the internal structure of an R object and generate summary statistics of the training dataset
- The Training dataset contains 160 variables and 19,622 records
- The Testing dataset contains 160 variables and 20 records

```
# Load the required r packages
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(randomForest)
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## margin
```

```
dfTrain <- read.csv("pml-training.csv", header = TRUE, na.strings=c("NA", "#DIV/0!", ""))
```

```
dfTest <- read.csv("pml-testing.csv", header = TRUE, na.strings=c("NA", "#DIV/0!", ""))
```

```
# Get variable names
```

```
names(dfTrain)
```

```
## [1] "X" "user_name"
```

```
## [3] "raw_timestamp_part_1" "raw_timestamp_part_2"
```

##	[5]	"cvtd_timestamp"	"new_window"
##	[7]	"num_window"	"roll_belt"
##	[9]	"pitch_belt"	"yaw_belt"
##	[11]	"total_accel_belt"	"kurtosis_roll_belt"
##	[13]	"kurtosis_picth_belt"	"kurtosis_yaw_belt"
##	[15]	"skewness_roll_belt"	"skewness_roll_belt.1"
##	[17]	"skewness_yaw_belt"	"max_roll_belt"
##	[19]	"max_picth_belt"	"max_yaw_belt"
##	[21]	"min_roll_belt"	"min_pitch_belt"
##	[23]	"min_yaw_belt"	"amplitude_roll_belt"
##	[25]	"amplitude_pitch_belt"	"amplitude_yaw_belt"
##	[27]	"var_total_accel_belt"	"avg_roll_belt"
##	[29]	"stddev_roll_belt"	"var_roll_belt"
##	[31]	"avg_pitch_belt"	"stddev_pitch_belt"
##	[33]	"var_pitch_belt"	"avg_yaw_belt"
##	[35]	"stddev_yaw_belt"	"var_yaw_belt"
##	[37]	"gyros_belt_x"	"gyros_belt_y"
##	[39]	"gyros_belt_z"	"accel_belt_x"
##	[41]	"accel_belt_y"	"accel_belt_z"
##	[43]	"magnet_belt_x"	"magnet_belt_y"
##	[45]	"magnet_belt_z"	"roll_arm"
##	[47]	"pitch_arm"	"yaw_arm"
##	[49]	"total_accel_arm"	"var_accel_arm"
##	[51]	"avg_roll_arm"	"stddev_roll_arm"
##	[53]	"var_roll_arm"	"avg_pitch_arm"
##	[55]	"stddev_pitch_arm"	"var_pitch_arm"
##	[57]	"avg_yaw_arm"	"stddev_yaw_arm"
##	[59]	"var_yaw_arm"	"gyros_arm_x"
##	[61]	"gyros_arm_y"	"gyros_arm_z"
##	[63]	"accel_arm_x"	"accel_arm_y"
##	[65]	"accel_arm_z"	"magnet_arm_x"
##	[67]	"magnet_arm_y"	"magnet_arm_z"
##	[69]	"kurtosis_roll_arm"	"kurtosis_picth_arm"
##	[71]	"kurtosis_yaw_arm"	"skewness_roll_arm"
##	[73]	"skewness_pitch_arm"	"skewness_yaw_arm"
##	[75]	"max_roll_arm"	"max_picth_arm"
##	[77]	"max_yaw_arm"	"min_roll_arm"
##	[79]	"min_pitch_arm"	"min_yaw_arm"
##	[81]	"amplitude_roll_arm"	"amplitude_pitch_arm"
##	[83]	"amplitude_yaw_arm"	"roll_dumbbell"
##	[85]	"pitch_dumbbell"	"yaw_dumbbell"
##	[87]	"kurtosis_roll_dumbbell"	"kurtosis_picth_dumbbell"
##	[89]	"kurtosis_yaw_dumbbell"	"skewness_roll_dumbbell"
##	[91]	"skewness_pitch_dumbbell"	"skewness_yaw_dumbbell"
##	[93]	"max_roll_dumbbell"	"max_picth_dumbbell"
##	[95]	"max_yaw_dumbbell"	"min_roll_dumbbell"
##	[97]	"min_pitch_dumbbell"	"min_yaw_dumbbell"
##	[99]	"amplitude_roll_dumbbell"	"amplitude_pitch_dumbbell"
##	[101]	"amplitude_yaw_dumbbell"	"total_accel_dumbbell"
##	[103]	"var_accel_dumbbell"	"avg_roll_dumbbell"
##	[105]	"stddev_roll_dumbbell"	"var_roll_dumbbell"
##	[107]	"avg_pitch_dumbbell"	"stddev_pitch_dumbbell"
##	[109]	"var_pitch_dumbbell"	"avg_yaw_dumbbell"
##	[111]	"stddev_yaw_dumbbell"	"var_yaw_dumbbell"

```
## [113] "gyros_dumbbell_x"      "gyros_dumbbell_y"
## [115] "gyros_dumbbell_z"      "accel_dumbbell_x"
## [117] "accel_dumbbell_y"      "accel_dumbbell_z"
## [119] "magnet_dumbbell_x"     "magnet_dumbbell_y"
## [121] "magnet_dumbbell_z"     "roll_forearm"
## [123] "pitch_forearm"         "yaw_forearm"
## [125] "kurtosis_roll_forearm" "kurtosis_pitch_forearm"
## [127] "kurtosis_yaw_forearm"  "skewness_roll_forearm"
## [129] "skewness_pitch_forearm" "skewness_yaw_forearm"
## [131] "max_roll_forearm"      "max_pitch_forearm"
## [133] "max_yaw_forearm"       "min_roll_forearm"
## [135] "min_pitch_forearm"     "min_yaw_forearm"
## [137] "amplitude_roll_forearm" "amplitude_pitch_forearm"
## [139] "amplitude_yaw_forearm" "total_accel_forearm"
## [141] "var_accel_forearm"     "avg_roll_forearm"
## [143] "stddev_roll_forearm"   "var_roll_forearm"
## [145] "avg_pitch_forearm"     "stddev_pitch_forearm"
## [147] "var_pitch_forearm"     "avg_yaw_forearm"
## [149] "stddev_yaw_forearm"    "var_yaw_forearm"
## [151] "gyros_forearm_x"       "gyros_forearm_y"
## [153] "gyros_forearm_z"       "accel_forearm_x"
## [155] "accel_forearm_y"       "accel_forearm_z"
## [157] "magnet_forearm_x"      "magnet_forearm_y"
## [159] "magnet_forearm_z"      "classe"
```

```
str(dfTrain)
```

```
## 'data.frame':   19622 obs. of  160 variables:
## $ X : int  1 2 3 4 5 6 7 8 9 10 ...
## $ user_name : Factor w/ 6 levels "adelmo","carlitos",...: 2 2 2 2 2 2 2 2 2 2 ...
## $ raw_timestamp_part_1 : int  1323084231 1323084231 1323084231 1323084232 1323084232 1323084232 ...
## $ raw_timestamp_part_2 : int  788290 808298 820366 120339 196328 304277 368296 440390 484323 484...
## $ cvtd_timestamp : Factor w/ 20 levels "02/12/2011 13:32",...: 9 9 9 9 9 9 9 9 9 9 ...
## $ new_window : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
## $ num_window : int  11 11 11 12 12 12 12 12 12 12 ...
## $ roll_belt : num  1.41 1.41 1.42 1.48 1.48 1.45 1.42 1.42 1.43 1.45 ...
## $ pitch_belt : num  8.07 8.07 8.07 8.05 8.07 8.06 8.09 8.13 8.16 8.17 ...
## $ yaw_belt : num  -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 ...
## $ total_accel_belt : int  3 3 3 3 3 3 3 3 3 3 ...
## $ kurtosis_roll_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_pitch_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_yaw_belt : logi  NA NA NA NA NA NA ...
## $ skewness_roll_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_roll_belt.1 : num  NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_yaw_belt : logi  NA NA NA NA NA NA ...
## $ max_roll_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_belt : int  NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ min_roll_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_belt : int  NA NA NA NA NA NA NA NA NA NA ...
## $ min_yaw_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_roll_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_pitch_belt : int  NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_yaw_belt : num  NA NA NA NA NA NA NA NA NA NA ...
## $ var_total_accel_belt : num  NA NA NA NA NA NA NA NA NA NA ...
```

```

## $ avg_roll_belt      : num NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_roll_belt   : num NA NA NA NA NA NA NA NA NA NA ...
## $ var_roll_belt      : num NA NA NA NA NA NA NA NA NA NA ...
## $ avg_pitch_belt     : num NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_pitch_belt  : num NA NA NA NA NA NA NA NA NA NA ...
## $ var_pitch_belt     : num NA NA NA NA NA NA NA NA NA NA ...
## $ avg_yaw_belt       : num NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_yaw_belt    : num NA NA NA NA NA NA NA NA NA NA ...
## $ var_yaw_belt       : num NA NA NA NA NA NA NA NA NA NA ...
## $ gyros_belt_x       : num 0 0.02 0 0.02 0.02 0.02 0.02 0.02 0.02 0.03 ...
## $ gyros_belt_y       : num 0 0 0 0 0.02 0 0 0 0 0 ...
## $ gyros_belt_z       : num -0.02 -0.02 -0.02 -0.03 -0.02 -0.02 -0.02 -0.02 -0.02 0 ...
## $ accel_belt_x       : int -21 -22 -20 -22 -21 -21 -22 -22 -20 -21 ...
## $ accel_belt_y       : int 4 4 5 3 2 4 3 4 2 4 ...
## $ accel_belt_z       : int 22 22 23 21 24 21 21 21 24 22 ...
## $ magnet_belt_x      : int -3 -7 -2 -6 -6 0 -4 -2 1 -3 ...
## $ magnet_belt_y      : int 599 608 600 604 600 603 599 603 602 609 ...
## $ magnet_belt_z      : int -313 -311 -305 -310 -302 -312 -311 -313 -312 -308 ...
## $ roll_arm           : num -128 -128 -128 -128 -128 -128 -128 -128 -128 -128 ...
## $ pitch_arm          : num 22.5 22.5 22.5 22.1 22.1 22 21.9 21.8 21.7 21.6 ...
## $ yaw_arm            : num -161 -161 -161 -161 -161 -161 -161 -161 -161 -161 ...
## $ total_accel_arm    : int 34 34 34 34 34 34 34 34 34 34 ...
## $ var_accel_arm      : num NA NA NA NA NA NA NA NA NA NA ...
## $ avg_roll_arm       : num NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_roll_arm    : num NA NA NA NA NA NA NA NA NA NA ...
## $ var_roll_arm       : num NA NA NA NA NA NA NA NA NA NA ...
## $ avg_pitch_arm      : num NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_pitch_arm   : num NA NA NA NA NA NA NA NA NA NA ...
## $ var_pitch_arm      : num NA NA NA NA NA NA NA NA NA NA ...
## $ avg_yaw_arm        : num NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_yaw_arm     : num NA NA NA NA NA NA NA NA NA NA ...
## $ var_yaw_arm        : num NA NA NA NA NA NA NA NA NA NA ...
## $ gyros_arm_x        : num 0 0.02 0.02 0.02 0 0.02 0 0.02 0.02 0.02 ...
## $ gyros_arm_y        : num 0 -0.02 -0.02 -0.03 -0.03 -0.03 -0.03 -0.02 -0.03 -0.03 ...
## $ gyros_arm_z        : num -0.02 -0.02 -0.02 0.02 0 0 0 0 -0.02 -0.02 ...
## $ accel_arm_x        : int -288 -290 -289 -289 -289 -289 -289 -289 -288 -288 ...
## $ accel_arm_y        : int 109 110 110 111 111 111 111 111 109 110 ...
## $ accel_arm_z        : int -123 -125 -126 -123 -123 -122 -125 -124 -122 -124 ...
## $ magnet_arm_x       : int -368 -369 -368 -372 -374 -369 -373 -372 -369 -376 ...
## $ magnet_arm_y       : int 337 337 344 344 337 342 336 338 341 334 ...
## $ magnet_arm_z       : int 516 513 513 512 506 513 509 510 518 516 ...
## $ kurtosis_roll_arm  : num NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_pitch_arm : num NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_yaw_arm   : num NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_roll_arm  : num NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_pitch_arm : num NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_yaw_arm   : num NA NA NA NA NA NA NA NA NA NA ...
## $ max_roll_arm       : num NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_arm      : num NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_arm        : int NA NA NA NA NA NA NA NA NA NA ...
## $ min_roll_arm       : num NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_arm      : num NA NA NA NA NA NA NA NA NA NA ...
## $ min_yaw_arm        : int NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_roll_arm : num NA NA NA NA NA NA NA NA NA NA ...

```

```
## $ amplitude_pitch_arm      : num  NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_yaw_arm        : int   NA NA NA NA NA NA NA NA NA NA NA ...
## $ roll_dumbbell            : num   13.1 13.1 12.9 13.4 13.4 ...
## $ pitch_dumbbell           : num   -70.5 -70.6 -70.3 -70.4 -70.4 ...
## $ yaw_dumbbell             : num   -84.9 -84.7 -85.1 -84.9 -84.9 ...
## $ kurtosis_roll_dumbbell    : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_pitch_dumbbell   : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_yaw_dumbbell     : logi   NA NA NA NA NA NA ...
## $ skewness_roll_dumbbell    : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_pitch_dumbbell   : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_yaw_dumbbell     : logi   NA NA NA NA NA NA ...
## $ max_roll_dumbbell        : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_dumbbell       : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_dumbbell         : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_roll_dumbbell        : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_dumbbell       : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_yaw_dumbbell         : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_roll_dumbbell   : num   NA NA NA NA NA NA NA NA NA NA NA ...
## [list output truncated]
```

```
dim(dfTest)
```

```
## [1] 20 160
```

```
#summary(dfTrain)
```

```
summary(dfTrain$classe)
```

```
##      A      B      C      D      E
## 5580 3797 3422 3216 3607
```

## Data Processing: Cleaning and Preparation

- Remove the first seven descriptive variables/fields (X/Id, user\_name, raw\_timestamp\_part\_1, raw\_timestamp\_part\_2, cvtd\_timestamp, new\_window, num\_window) from both data sets that will not help predict the manner in which an exercise was completed.
- Remove the variables/fields from the data set that contain missing values
- Remove Near Zero Variance Variables
- The resulting Training and Testing datasets both have 53 variables/fields the last of which is the classe variable/field
- Split the cleaned training data set into a training set (75%) that will be used for prediction and a testing/validation set (25%) that will be used to determine out-of-sample errors

```
dfTrain <- dfTrain[, -c(1:7)]
```

```
dfTest <- dfTest[, -c(1:7)]
```

```
dfTrain <- dfTrain[, colSums(is.na(dfTrain)) == 0]
```

```
dfTest <- dfTest[, colSums(is.na(dfTest)) == 0]
```

```
#Check if there are Near Zero Variance Variables to remove
```

```
nzVar <- nearZeroVar(dfTrain, saveMetrics = TRUE)
```

```
nzVar
```

```
##               freqRatio percentUnique zeroVar  nzv
## roll_belt      1.101904      6.7781062  FALSE FALSE
## pitch_belt     1.036082      9.3772296  FALSE FALSE
```

## yaw_belt	1.058480	9.9734991	FALSE	FALSE
## total_accel_belt	1.063160	0.1477933	FALSE	FALSE
## gyros_belt_x	1.058651	0.7134849	FALSE	FALSE
## gyros_belt_y	1.144000	0.3516461	FALSE	FALSE
## gyros_belt_z	1.066214	0.8612782	FALSE	FALSE
## accel_belt_x	1.055412	0.8357966	FALSE	FALSE
## accel_belt_y	1.113725	0.7287738	FALSE	FALSE
## accel_belt_z	1.078767	1.5237998	FALSE	FALSE
## magnet_belt_x	1.090141	1.6664968	FALSE	FALSE
## magnet_belt_y	1.099688	1.5187035	FALSE	FALSE
## magnet_belt_z	1.006369	2.3290184	FALSE	FALSE
## roll_arm	52.338462	13.5256345	FALSE	FALSE
## pitch_arm	87.256410	15.7323412	FALSE	FALSE
## yaw_arm	33.029126	14.6570176	FALSE	FALSE
## total_accel_arm	1.024526	0.3363572	FALSE	FALSE
## gyros_arm_x	1.015504	3.2769341	FALSE	FALSE
## gyros_arm_y	1.454369	1.9162165	FALSE	FALSE
## gyros_arm_z	1.110687	1.2638875	FALSE	FALSE
## accel_arm_x	1.017341	3.9598410	FALSE	FALSE
## accel_arm_y	1.140187	2.7367241	FALSE	FALSE
## accel_arm_z	1.128000	4.0362858	FALSE	FALSE
## magnet_arm_x	1.000000	6.8239731	FALSE	FALSE
## magnet_arm_y	1.056818	4.4439914	FALSE	FALSE
## magnet_arm_z	1.036364	6.4468454	FALSE	FALSE
## roll_dumbbell	1.022388	84.2065029	FALSE	FALSE
## pitch_dumbbell	2.277372	81.7449801	FALSE	FALSE
## yaw_dumbbell	1.132231	83.4828254	FALSE	FALSE
## total_accel_dumbbell	1.072634	0.2191418	FALSE	FALSE
## gyros_dumbbell_x	1.003268	1.2282132	FALSE	FALSE
## gyros_dumbbell_y	1.264957	1.4167771	FALSE	FALSE
## gyros_dumbbell_z	1.060100	1.0498420	FALSE	FALSE
## accel_dumbbell_x	1.018018	2.1659362	FALSE	FALSE
## accel_dumbbell_y	1.053061	2.3748853	FALSE	FALSE
## accel_dumbbell_z	1.133333	2.0894914	FALSE	FALSE
## magnet_dumbbell_x	1.098266	5.7486495	FALSE	FALSE
## magnet_dumbbell_y	1.197740	4.3012945	FALSE	FALSE
## magnet_dumbbell_z	1.020833	3.4451126	FALSE	FALSE
## roll_forearm	11.589286	11.0895933	FALSE	FALSE
## pitch_forearm	65.983051	14.8557741	FALSE	FALSE
## yaw_forearm	15.322835	10.1467740	FALSE	FALSE
## total_accel_forearm	1.128928	0.3567424	FALSE	FALSE
## gyros_forearm_x	1.059273	1.5187035	FALSE	FALSE
## gyros_forearm_y	1.036554	3.7763735	FALSE	FALSE
## gyros_forearm_z	1.122917	1.5645704	FALSE	FALSE
## accel_forearm_x	1.126437	4.0464784	FALSE	FALSE
## accel_forearm_y	1.059406	5.1116094	FALSE	FALSE
## accel_forearm_z	1.006250	2.9558659	FALSE	FALSE
## magnet_forearm_x	1.012346	7.7667924	FALSE	FALSE
## magnet_forearm_y	1.246914	9.5403119	FALSE	FALSE
## magnet_forearm_z	1.000000	8.5771073	FALSE	FALSE
## classe	1.469581	0.0254816	FALSE	FALSE

```
#dim(nzVar)
#head(nzVar, 60)
```

```
dfTrain <- dfTrain[, !nzVar$nzv]
dfTest <- dfTest[, !nzVar$nzv]
dim(dfTrain)

## [1] 19622    53

dfInTrain <- createDataPartition(dfTrain$classe, p = 0.75, list = FALSE)
dfPredict <- dfTrain[dfInTrain, ]
dfValidate <- dfTrain[-dfInTrain, ]
```

## Model Fitting

- set.seed for pseudo-random number generation and ensure reproducible results
- A predictive model is fitted to predict the manner/class type in which an exercise was completed using Random Forest algorithm
- Random Forest algorithm is selected here because it is one of the most accurate learning algorithms available and produces highly accurate classifier for many datasets. It provides estimates of what variables are important in the classification and handles correlated covariates & outliers.
- A 5-fold cross validation (cv) resampling method is applied to the algorithm
- The results are predicted using the validation data set
- The results are compared using a confusionMatrix: a cross-tabulation of observed and predicted classes with associated statistics.
- The accuracy/overall agreement rate and Kappa are computed

```
set.seed(25)
fitControl <- trainControl(method='cv', number = 5)
modFitRf <- train(classe ~ ., data = dfPredict, method = "rf", trControl = fitControl)
#print(modFitRf)
modFitRf
```

```
## Random Forest
##
## 14718 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 11775, 11773, 11775, 11775, 11774
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa
##    2    0.9910315 0.9886536
##   27    0.9915748 0.9893420
##   52    0.9843724 0.9802286
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

```
predictRf <- predict(modFitRf, dfValidate)

confusionMatrix(dfValidate$classe, predictRf)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1394    0    0    0    1
##           B   10  937    1    1    0
##           C    0    7  843    5    0
##           D    0    2   13  789    0
##           E    0    0    0    4  897
##
## Overall Statistics
##
##           Accuracy : 0.991
##           95% CI : (0.988, 0.9935)
##           No Information Rate : 0.2863
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9886
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9929  0.9905  0.9837  0.9875  0.9989
## Specificity      0.9997  0.9970  0.9970  0.9963  0.9990
## Pos Pred Value   0.9993  0.9874  0.9860  0.9813  0.9956
## Neg Pred Value   0.9972  0.9977  0.9965  0.9976  0.9998
## Prevalence       0.2863  0.1929  0.1748  0.1629  0.1831
## Detection Rate   0.2843  0.1911  0.1719  0.1609  0.1829
## Detection Prevalence 0.2845  0.1935  0.1743  0.1639  0.1837
## Balanced Accuracy 0.9963  0.9937  0.9903  0.9919  0.9989
accuracy <- postResample(predictRf, dfValidate$classe)
accuracy

## Accuracy      Kappa
## 0.9910277 0.9886485
```

## Conclusions & Test Data Set Prediction

- The Random Forest algorithm performed well with an accuracy of 0.995. The expected out-of-sample error rate is estimated at 0.005 (1 - accuracy).
- Therefore, the Random Forest predictive model is applied to the 20 test cases available in the test data set. We can expect that few of the test samples will be misclassified based on the accurate shown on the cross-validation data set.
- 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
- B A B A A E D B A A B C B A E E A B B B

```
predictRf <- predict(modFitRf, dfTest)
predictRf

## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```